


Docker 101 Tutorial

2022-09-29

Agenda

- Virtualization: Virtual Machine and Container
- Docker Architecture and Principles
- Docker Use Cases
- Podman: Daemonless and Rootless Container Engine


Physical Server Architecture



Physical Server: Problems


- Limited servers
- Run old and incompatible software
- Develop cross-platform applications
- Try new operating systems
- Sandbox environment, such as virus testing

Virtual Machine: Type-2 Hypervisor



- Hypervisor: Virtual Machine Manager/Monitor(VMM)
- Hosted Hypervisor
- Commercial products in 2000s
- Typical Type-2 Hypervisor Products
 - VMWare Workstation Player/Pro
 - Oracle VM VirtualBox
 - QEMU

Virtual Machine: Type-1 Hypervisor



- Bare metal hypervisor
- Typical Type-1 Hypervisor Products
 - VMWare ESX/ESXi(vSphere)
 - Microsoft Hyper-V
 - Oracle VM Server
 - Citrix Hypervisor(XenServer)
 - KVM


Type-1 Virtual Machine: Real-World Examples

- AWS Elastic Compute Cloud(EC2)
- Microsoft Azure Virtual Machine(Azure)
- Google Cloud Platform(GCP)
- Alibaba Cloud Elastic Compute Service(ECS)
- Tencent Cloud Virtual Machine(CVM)
- SmartX(native hypervisor ELF, KVM based)

Virtual Machine: Problems

- Resource intensive
- Slow startup and shutdown
- Deployment complexity
- Scalability

Container

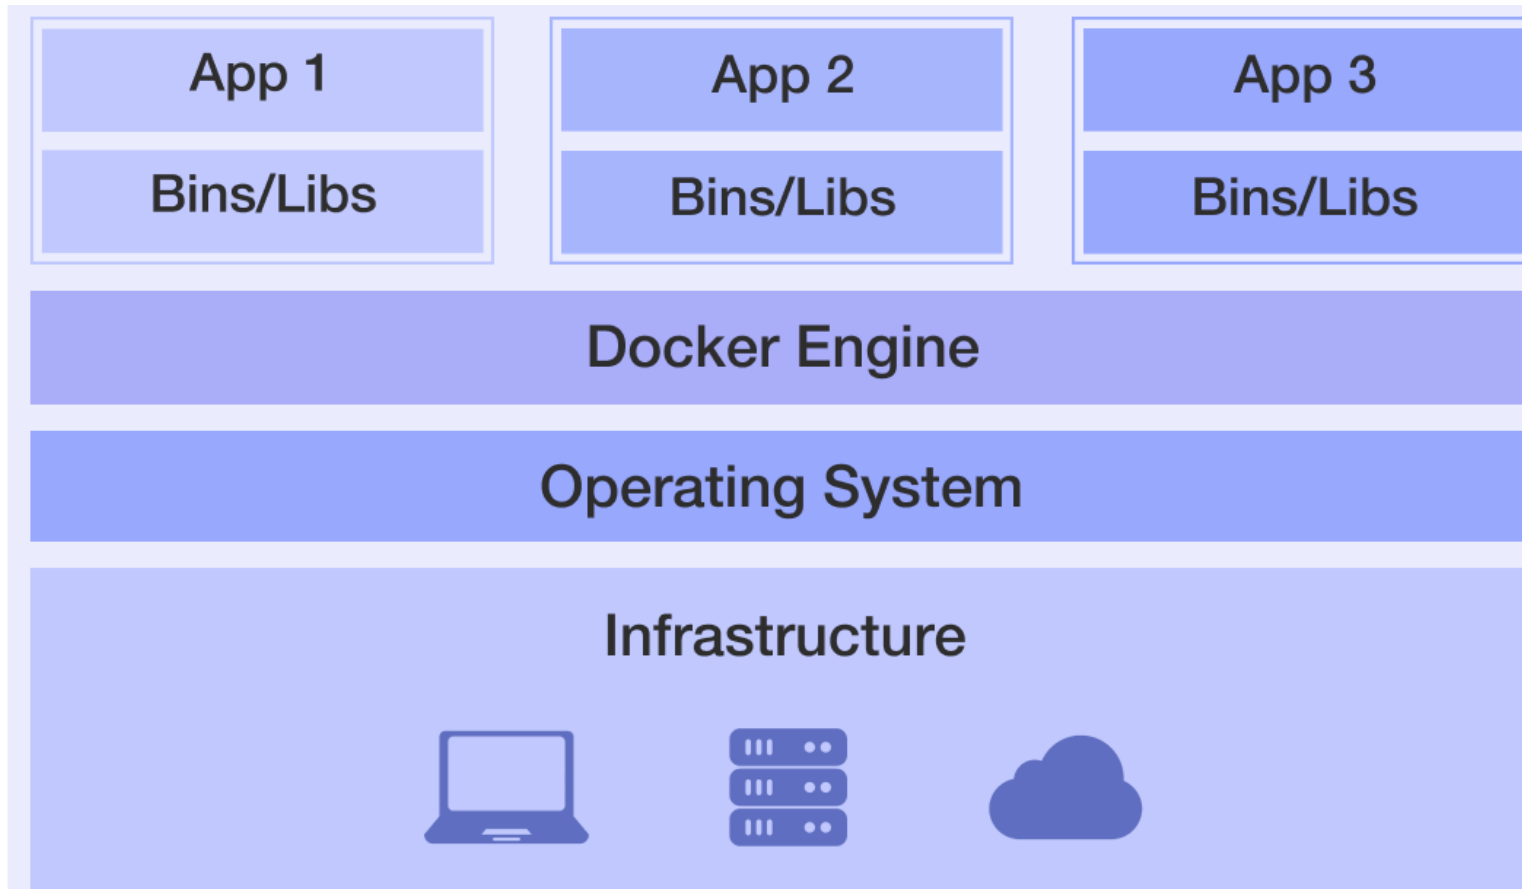


- OS-level virtualization, share OS kernel
- Isolated process
- LXC: Linux Container(2008)
- Namespaces and CGroups

Virtual Machine vs. Container

VMS	CONTAINERS
Heavyweight.	Lightweight.
Limited performance.	Native performance.
Each VM runs in its own OS.	All containers share the host OS.
Hardware-level virtualization.	OS virtualization.
Startup time in minutes.	Startup time in milliseconds.
Allocates required memory.	Requires less memory space.
Fully isolated and hence more secure.	Process-level isolation, possibly less secure.

Docker Infrastructure

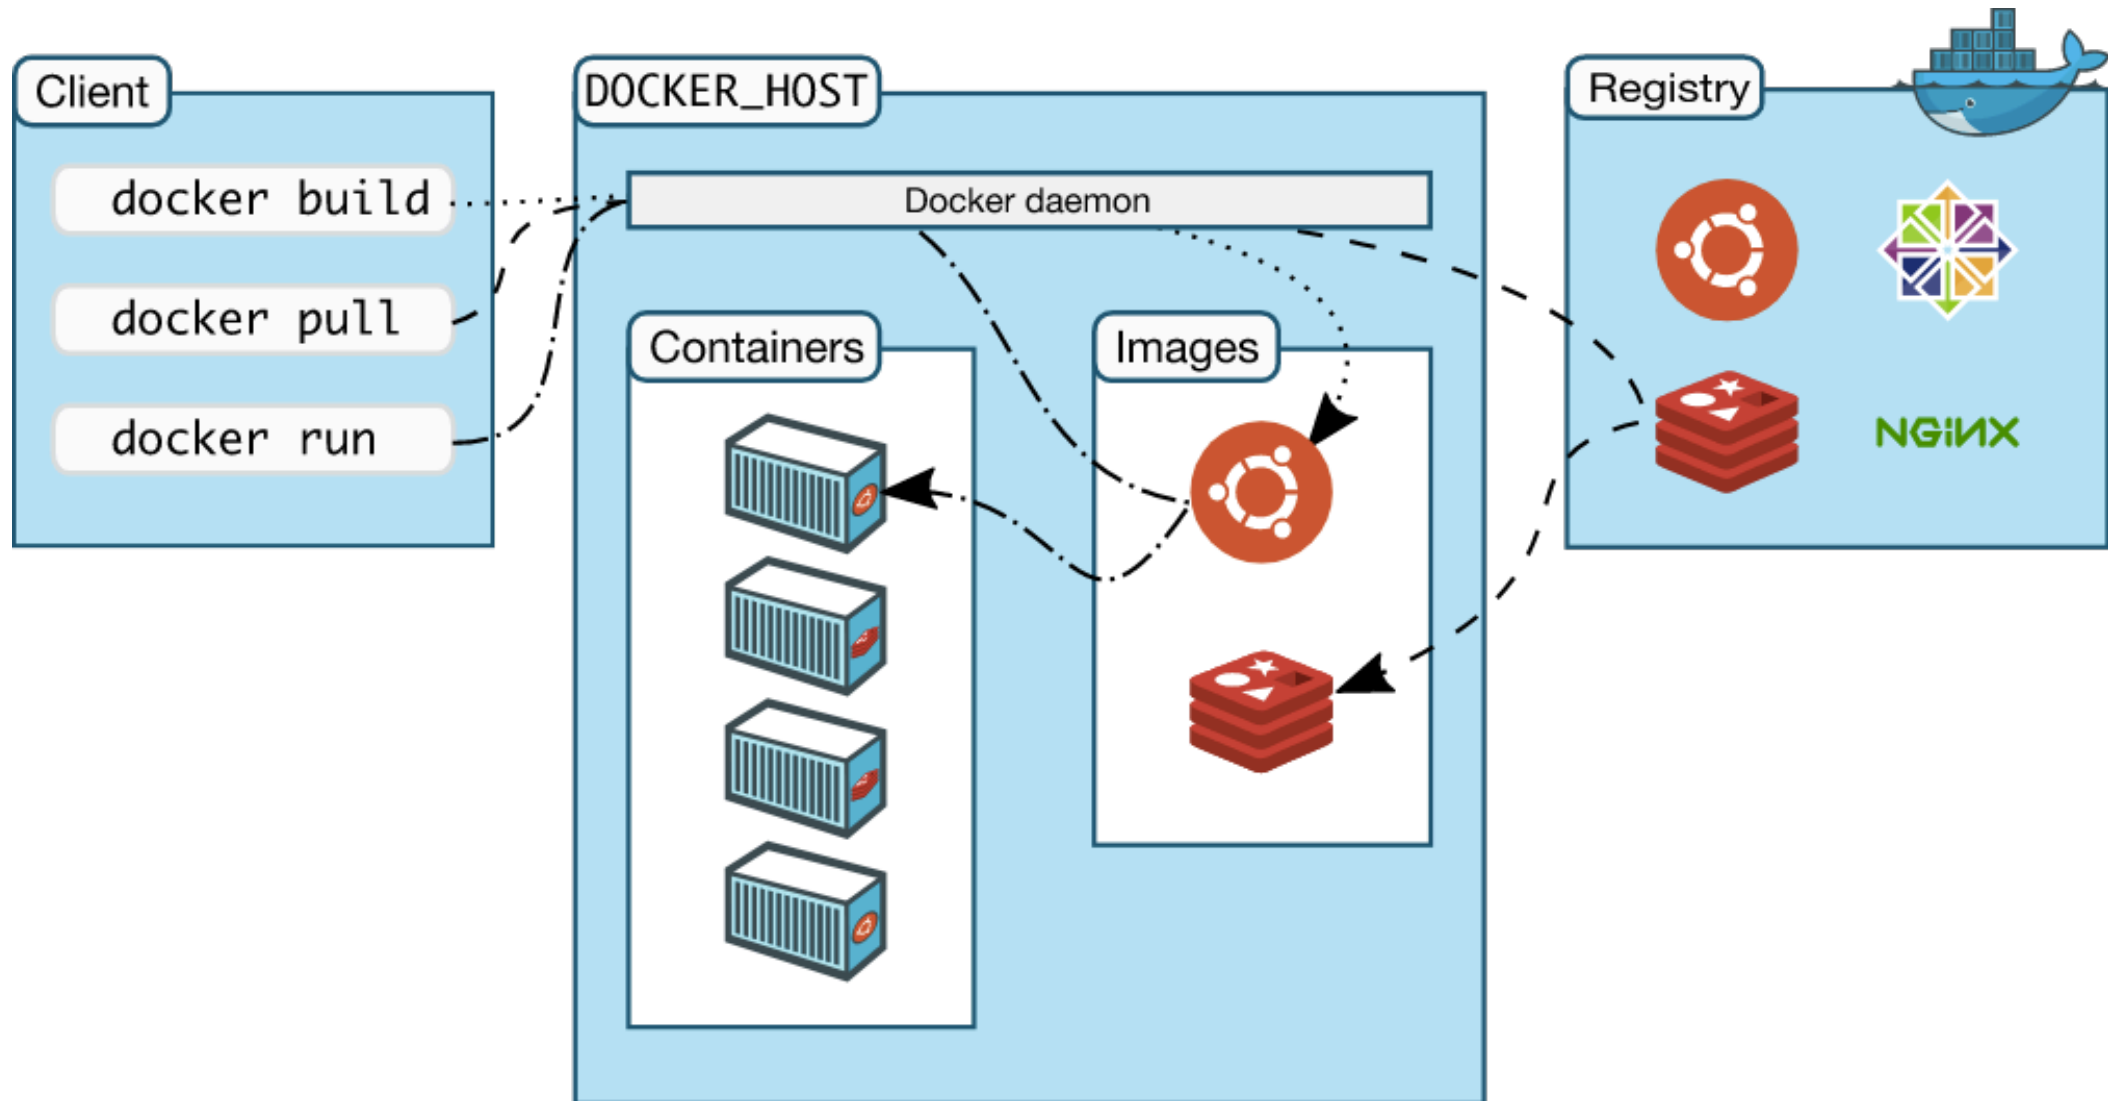


- Container technology based
- OS-level virtualization
- Innovation on LXC(2013)
- Docker Engine = dockerd + containerd + runC

Docker Added to LXC

- Portable deployment across machines
- Application-centric
- Automatic build: Dockerfile
- Versioning
- Component re-use
- Sharing
- Tool ecosystem

Docker Architecture: client/server



Docker Engine

- **Server:** It is the docker daemon called **dockerd**. It can create and manage docker images, containers, networks, etc.
- **Rest API:** It is used to instruct docker daemon what to do.
- **Command Line Interface (CLI):** It is a client which is used to enter [docker commands](#).

Docker Registry

- Similar to git repository
- Store docker images
- Public registry: Docker Hub
- Private registry

Docker Objects

- images
- containers
- volumes
- networks

Docker Images

- read-only template with instructions for creating a Docker container
- image = base image + customization
- package: code+configuration+dependencies
- Dockerfile -> docker image

Dockerfile

```
FROM python

WORKDIR /home/docker/python-docker-demo

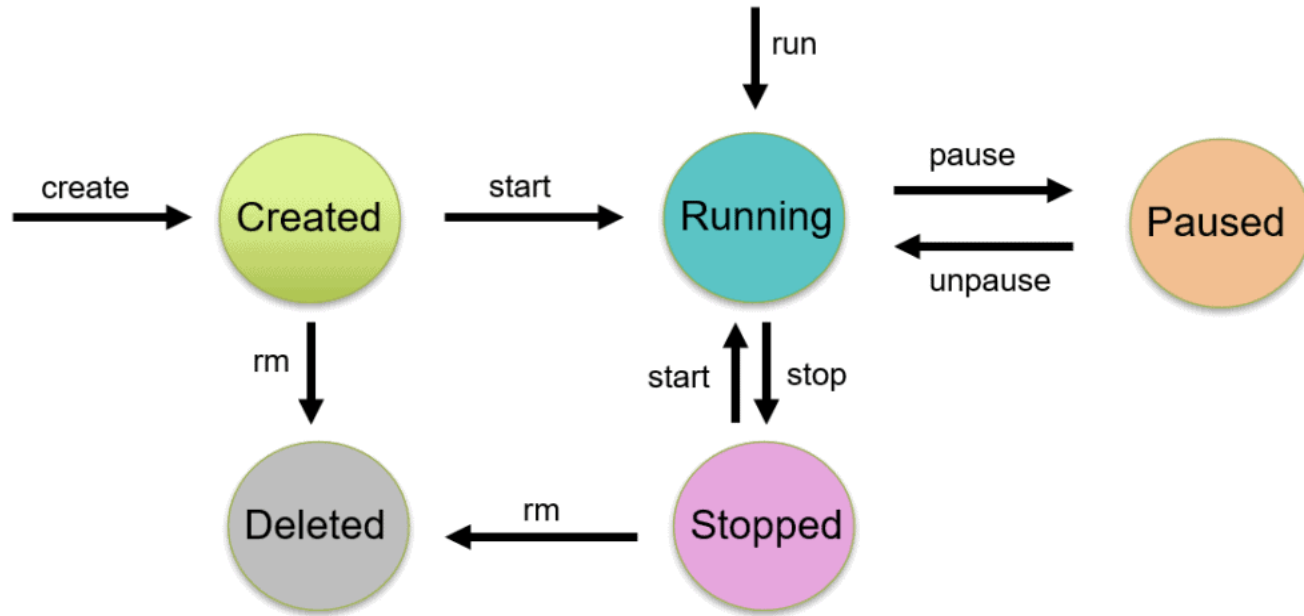
COPY . .

RUN pip3 install -r requirements.txt -i https://pypi.tuna.tsinghua.edu.cn/simple

ENTRYPOINT ["python", "main.py"]
```

docker build -t python-docker-image:tag path/to/Dockerfile
<https://docs.docker.com/engine/reference/builder/>

Docker Containers



- container: runnable instance of image
- All the applications and their environment run inside the container
- Docker API or CLI to operate container

Volumes

- Store the persisting data generated by docker and used by Docker containers.
- Volume's content exists outside the lifecycle of a container.

Networks

- **Bridge:** It is the default network driver for a container.
- **Host:** no network isolation between host and container.
- **Overlay:** This network enables swarm services to communicate with each other.
- **None:** disable all the networking.
- **macvlan:** Assigns mac address to containers to make them look like physical devices.

Docker Practice

- Docker installation
- Docker use cases
- Docker problems

Docker Installation

- Desktop
 - Docker Desktop
 - Linux, Mac, Windows(10)
 - <https://docs.docker.com/desktop/>
- Server
 - CentOS, Debian, Fedora, Ubuntu, RHEL, etc.
 - <https://docs.docker.com/engine/install/>

Docker Commands

- management commands
- commands
- docker --help
- <https://docs.docker.com/engine/reference/commandline/docker/>

Docker Scenarios

- Application isolation
- Build portable environment
- Microservices
- CI/CD

Docker Problems

- Root privilege: daemon binds unix socket
- Docker user group
- Security problems
- Rootless mode: Docker Engine v20.10, limitations
- <https://docs.docker.com/engine/release-notes/>
- <https://docs.docker.com/engine/security/rootless/>

Podman: Pod Manager

- Daemonless and rootless container engine
- Docker command compatible
- **alias docker=podman**
- <https://podman.io/>

Key Takeaways

- Virtualization Techniques: VM and Container
- VM: type-1 hypervisor vs. type-2 hypervisor
- Docker: A Container Engine
- Docker can't replace VM
- Package up application and all its dependencies
- Podman: rootless container engine

Thoughts

- Multiple containers on a single host machine?
 - Docker Compose
- Multiple containers across multiple host machines?
 - Docker Swarm
 - Kubernetes(K8s)

References

- <https://www.docker.com/>
- <https://podman.io/>
- https://yeasy.gitbook.io/docker_practice/
- <http://icyfenix.cn/immutable-infrastructure/container/>

Thanks
Q&A